

# DAIWANG KHERA

+91-9205331936 | daiwangk@gmail.com | Gurgaon, India | Portfolio | LinkedIn | GitHub

## EDUCATION

**Dronacharya College of Engineering**  
*B.Tech in Computer Science | CGPA: 8.0/10*

Gurugram, Haryana  
*Expected 2026*

## TECHNICAL SKILLS

**Programming:** Python, JavaScript, SQL  
**AI/ML:** RAG, Corrective RAG (CRAG), Prompt Engineering, LLM Integration (Groq, Ollama, LLaMA), QLoRA  
Fine-Tuning, Embeddings, LangGraph, LangChain, OpenCV  
**Frameworks:** FastAPI, Flask, Rasa, Streamlit, LlamaIndex, LlamaParse, ChromaDB  
**Databases:** SQLite, ChromaDB (Vector DB)  
**DevOps & Tools:** Docker, Git, Linux, n8n, Ragas, Pydantic v2

## EXPERIENCE

**AI/ML Intern** Jul 2025 – Aug 2025  
*Black Box Limited*

- Built a conversational retail assistant using Rasa and FastAPI for product discovery and order tracking
- Integrated Llama 3.2 via Ollama for product summarization, reducing support queries by ~30%
- Fine-tuned Gemma-1B with QLoRA, improving intent recognition accuracy from 74% to 82%
- Designed prompt templates and coordinated demo-ready AI workflows with internal teams

**Frontend Development Intern** Nov 2024 – Jan 2025  
*Medecro Technologies Pvt Ltd*

- Built scheduling and inventory dashboard components serving the internal operations team
- Implemented responsive layouts and integrated REST APIs using Git-based workflows
- Collaborated with cross-functional development teams via strict version control practices

**Machine Learning Intern** Jun 2024 – Jul 2024  
*Z Plus Surveillance Security Pvt Ltd*

- Built motion detection pipeline using OpenCV with blur filtering for surveillance video feeds

## PROJECTS

**Legal & Financial Document Intelligence** | *LangGraph, FastAPI, ChromaDB, Groq, Streamlit* | Demo 2026

- Engineered a Corrective RAG pipeline using LangGraph that grades retrieved chunks and rewrites queries automatically, eliminating hallucinations
- Implemented hierarchical chunking via LlamaIndex, preserving clause-to-section context in legal documents
- Designed dual-LLM architecture: Llama 3.1-8B for grading + Llama 3.3-70B for generation, optimizing latency and accuracy
- Deployed full-stack app (FastAPI + Streamlit) with Pydantic v2 validation and Ragas evaluation on Hugging Face Spaces

**Clothing & Retail AI Assistant** | *Rasa, FastAPI, SQLite, Docker* | GitHub Internship

- Enabled product discovery, order tracking, and recommendations via conversational interface
- Integrated Llama 3.2 summarization to reduce manual catalog lookups by ~25%

**Real-Time Motion Detection** | *Python, OpenCV* 2024

- Built surveillance pipeline using frame differencing and contour detection at ~45ms/frame latency

**Automated YouTube Shorts Pipeline** | *n8n, Python, Groq API, LLaMA 3.3* Side Project

- Built end-to-end workflow: Reddit ingestion, LLaMA 3.3 narration, video generation, and scheduled YouTube uploads
- Automated metadata logging via Google Sheets and content ranking using n8n workflows

## ACHIEVEMENTS & CERTIFICATIONS

AIR 18 – NCAT 2023 | Finalist – Medecro HealthHack 2024 | NPTEL Python Programming

**Leadership:** Led team during AI bootcamp project, coordinating solution design and presentation